

Whose Intelligence?

Building Agentic AI in Africa without Frontier Model Dependency

Session Format: Forums & Dialogues **Session Length:** 1.5h **Affiliation:** ML Collective

Abstract

The promise of Agentic AI is seductive: autonomous systems that perceive, plan, and act, thereby reducing the friction between human intent and outcome. However, Agentic AI is arriving in Africa on someone else's terms. Building agentic systems today almost inevitably means building on frontier models trained predominantly on Western data, optimized for Western languages, and aligned to Western values; yet deployed to negotiate, decide, and act on behalf of African users. Every API call, every default behavior, every quietly embedded assumption carries a worldview that was not built with Africa in mind.

This session takes a positional stance: the current default is insufficient and, in some cases, harmful. Through a startup pitch, positional talks, and a structured fishbowl dialogue, we pressure-test four tensions (cultural alignment, infrastructure dependency, value extraction, and the capability trap) working toward a community-driven position on sovereign agentic AI for Africa. The conversation will not end in the room; outputs will be compiled into an open preprint.

Goals & Objectives

This session aims to:

- Critically examine the assumptions embedded in building Agentic AI on top of frontier models in the African context.
- Move the African Agentic AI community beyond adoption conversations toward ownership and sovereignty conversations.
- Surface honest, grounded perspectives on the trade-offs between feasibility and dependency when building agentic systems under resource constraints.
- Contribute a community-driven position on what responsible, contextually-grounded Agentic AI development looks like for Africa.

Expected Outcomes for Participants

- A clearer framework for evaluating when and how to use frontier models in agentic AI tasks versus investing in local capacity.
- Exposure to concrete examples of agentic AI built for African realities, grounding the theoretical tensions in practice.
- A structured post-session report detailing key positions, points of divergence, and resulting recommendations, published as an open preprint.

Success Metrics

- Diversity and depth of perspectives represented across researchers, practitioners, and policy-makers.
- Quality of discussion generated around each of the four tensions, assessed through facilitator notes and participant feedback.
- Uptake of the session's outputs within the African AI community after the event.

Target Audience

This session is designed for practitioners, researchers, and policymakers at intermediate to advanced career stages who are building, studying, or shaping agentic AI systems in African contexts. Participants are expected to have familiarity with LLMs and agentic system concepts; this is not an introductory session. Startup founders navigating real infrastructure constraints will find particular value in the startup pitch and fishbowl segments.

Session Format & Justification

This session adopts a **forums and dialogues** format; specifically a hybrid of positional talks, a startup pitch, and a structured fishbowl dialogue. This format is deliberately chosen over a conventional workshop because the questions this session raises are not primarily pedagogical. They are contested, political, and unresolved. A lecture-style format would flatten the very tensions we are trying to surface. The fishbowl structure ensures that multiple perspectives (e.g., researcher, practitioner, policymaker, builder) are in active conversation rather than talking past each other from a panel.

Organising Team

Name	Affiliation	Role
Abraham Owodunni	The Ohio State University / ML Collective & Masakhane	Primary Contact, Lead Organizer
Chisom Chibuike	ML Collective	Primary Contact, Lead Organizer
Honour-Jesus Bezaleel	ML Collective	Organizer
Ifihanagbara Olushey	University of the Witwatersrand / ML Collective	Organizer
Steven Kolawole	Carnegie Mellon University / ML Collective	Organizer, Closing Remarks
Mardiyyah Oduwale	ML Collective	Organizer
Busayo Awobade	Intron Health / ML Collective	Organizer
Fatima Adebajo	ML Collective	Organizer
Rosanne Liu	Google DeepMind / ML Collective	Advisor, Speaker
Jason Yosinski	OpenAI / ML Collective	Advisor, Speaker

The organizing team spans academia (CMU, OSU, Wits), industry (Google DeepMind, OpenAI, Intron Health), and community organizations (ML Collective, Masakhane), with geographic representation across Nigeria, South Africa, and the United States. Team members have prior experience leading open-science ML communities (from grassroots initiatives like Data Science Nigeria, Google Developer Student Clubs, ML in Lagos to global programs like ML Collective) and organizing workshops at venues including NeurIPS, ICLR, and previous Indaba editions.

Speakers

Confirmed

- Chris Emezue (Lanfrica / Mila)
- Jason Yosinski (OpenAI / ML Collective)
- Rosanne Liu (Google DeepMind / ML Collective)

Invited (Pending Confirmation)

- Jade Abbott (Lelapa AI)
- Tobi Olatunji (Intron AI)
- Benjamin Rosman (University of the Witwatersrand)

Tentative Agenda

Time	Segment	Speaker(s)	Description
0:00 – 0:10	Opening Remarks	Session Organizers	Introduction of session framing and agenda
0:10 – 0:20	Opening Teaser	Chris Emezue	Provocation: Is Africa destined to be a consumer of global intelligence or a builder of its own sovereign AI systems?
0:20 – 0:35	African Startup Pitch	Tobi Olatunji	A fast-paced pitch showcasing a real agentic AI system built for African realities (e.g., education, agriculture, language tech). Grounds the conversation in practice and exposes gaps in infrastructure, data, and deployment. Closes with: “ <i>Why are we still dependent?</i> ”
0:35 – 0:55	Positional Talks	Rosanne Liu, Jason Yosinski, Jade Abbott/Ben Rosman, Chris Emezue	Four short positional talks (3–5 minutes each) on the four tensions: cultural alignment, infrastructure dependency, value extraction, and the capability trap.
0:55 – 1:20	Fishbowl Dialogue	All speakers + audience	Structured multi-perspective exchange. An inner circle of speakers engages the tensions directly; outer circle participants rotate in. Facilitator guides without leading.
1:20 – 1:30	Q&A + Closing Remarks	Steven Kola-wole, Abraham Owodunni	Open floor, synthesis of key insights, and clear calls to action: research directions, community collaboration, and building sovereign AI systems in Africa.

Interactivity & Engagement

The session is designed to be participatory from the first minute to the last:

- **Opening Teaser:** A live provocation designed to immediately surface the tensions and set the tone for honest, critical dialogue.
- **Startup Pitch:** Grounds abstract tensions in a concrete, lived example of agentic AI built for African realities.

- **Positional Talks:** Short, provocative framings of each tension, designed to create productive disagreement rather than consensus.
- **Fishbowl Dialogue:** The core engagement mechanism; participants are not an audience but active contributors to a structured exchange.
- **Live Position Capture:** Key positions and disagreements are documented throughout, feeding directly into the post-session preprint.

Diversity & Inclusion

The organizing team includes members from academia and industry with diverse geographic, gender, and career stage backgrounds across Nigeria, South Africa, and the United States. Speakers are drawn from African AI research, agentic systems development, and AI governance, providing a grounded, multi-dimensional view of the challenges at hand. Our session intentionally creates space for voices from community-driven research ecosystems — including Masakhane, ML Collective, and Cohere for AI — alongside those from frontier research institutions.

Sponsorship & Promotion

ML Collective has sponsored over 60 participants to attend the Indaba in previous years and plans to do the same in 2026, directly increasing workshop attendance. We will promote the session through:

- Social media (Twitter: 14.5K followers; LinkedIn: 1,006 followers)
- MLC mailing lists (4K subscribers)
- Community channels: AI Saturday Lagos, Cohere for AI Discord, ML Collective Discord
- In-person events: IndabaX, Data Science Africa

References

1. Toward an African Agenda for AI Safety (2025). <https://arxiv.org/pdf/2508.13179>
2. How Agentic AI Is Transforming African Banking — Oliver Wyman (2025). <https://www.oliverwyman.com/our-expertise/insights/2025/apr/agentic-ai-transforming-african-banking.html>
3. Decolonizing African NLP: A Survey on Power Dynamics and Data Colonialism — AfricaNLP @ ICLR 2024. <https://openreview.net/pdf?id=D698BEfwgv>
4. Decolonizing LLMs: An Ethnographic Framework for AI in African Contexts (2025). <https://www.researchgate.net/publication/388916215>
5. Toward a Decolonial Framework for Communicating AI in African Public Service — Ooko (2025). <https://journals.sagepub.com/doi/10.1177/27523543251374174>
6. AI Investment in Resource-Constrained African Economies — MDPI (2025). <https://www.mdpi.com/2673-4060/6/2/70>
7. Foresight Africa 2025–2030: AI and Emerging Technologies — Brookings. <https://www.brookings.edu/wp-content/uploads/2025/01/Foresight-Africa-2025-2030-Chapter-5.pdf>
8. Masakhane — community position on data sovereignty. <https://www.masakhane.io>

Appendix: The Four Tensions — Extended Framing

The promise of Agentic AI is seductive: autonomous systems that perceive, plan, and act, thereby reducing the friction between human intent and outcome. But as Africa stands at the threshold of this next wave, a harder question demands our attention before the enthusiasm carries us too far. Whose intelligence are these agents acting on?

Today, building agentic AI almost inevitably means building on top of frontier models like GPT, Gemini, Claude, and their successors. These systems were trained predominantly on Western data, optimized for Western languages, and aligned to Western values and assumptions about how the world works, what a “task” is, and what a “good outcome” looks like. When we deploy agents whose foundational models rely on less African data but negotiate and decide on behalf of African users, we are not just plugging into an API. We are embedding a particular worldview (one that does not entirely capture our nuance and context) into the infrastructure of African lives.

This session takes a positional stance: the current default approach to building Agentic AI solely around frontier models is insufficient in Africa, and in some cases harmful. We are not arguing against using frontier models. We are arguing that we have not been honest enough about what we are trading away when we do.

Tension 1: The Cultural Alignment Problem

Frontier models are not culturally neutral, and the gaps they carry across language, context, and epistemic tradition do not disappear when you fine-tune on a small Yoruba or Amharic dataset. The cultural assumptions baked into these systems (i.e., about what constitutes a valid question, a good answer, a trustworthy source) reflect the communities that produced the training data and shaped the RLHF feedback. Whose intelligence are we building on?

Tension 2: The Infrastructure Paradox

Africa faces real compute and data constraints, and the most feasible path to Agentic AI currently leads straight to dependency on the very systems that do not fully capture our nuances and context. The infrastructure required to train frontier-scale models from scratch is unavailable to most African institutions. Yet the alternative (permanent reliance on external APIs) carries its own risks. How do we pursue agentic capability in a constrained setting without deepening dependency on systems built elsewhere?

Tension 3: The Value Extraction Loop

When African developers and users interact with frontier models, they generate data and feedback that flows back to improve those models — for the benefit of labs elsewhere. We are, in effect, subsidizing the refinement of systems that were not built for us. The more African users adopt these systems, the more valuable they become to their creators, while the structural gap between African AI capacity and frontier AI capacity widens.

Tension 4: The Capability Trap

Every time we default to a frontier model API instead of investing in local capacity, we make it harder to ever not need that API. The shortcut today becomes the ceiling tomorrow. Agentic AI built solely on frontier model dependency decides, quietly, that Africa will always be a consumer of intelligence rather than a builder of it. Does this make or mar Africa’s sovereign intelligence?

This session does not pretend to resolve these tensions. It convenes researchers, practitioners, and policymakers to pressure-test them; to argue openly, rigorously, and without deference to the hype. The question is not whether Africa should build Agentic AI. The question is whether we are being clear-eyed about what we are building, for whom, and on whose terms.